

Connection Information:

This is the information we used when connecting to the database through python

```
cnx = connection.MySQLConnection(  
    user="hello",  
    password='1234',  
    host='34.134.85.82',  
    database='uiucbuses'  
)
```

team071database is the name of the instance

```
berniezhu573@cloudshell:~ (friendly-path-342117)$ gcloud sql connect team071database --user=root --quiet
```

Data Definition Language (DDL) Commands:

```
CREATE TABLE BusLine (  
    LineID VARCHAR(40) NOT NULL,  
    Color VARCHAR(10),  
    Number INTEGER,  
    PRIMARY KEY(LineID)  
);
```

```
CREATE TABLE BusStop (  
    StopID VARCHAR(20) NOT NULL,  
    StopLatitude REAL,  
    StopLongitude REAL,  
    PRIMARY KEY(StopID)  
);
```

```
CREATE TABLE Buses (  
    BusID VARCHAR(5) NOT NULL,  
    LineID VARCHAR(40),  
    StopOrigin VARCHAR(20),  
    StopDestination VARCHAR(20),  
    Direction VARCHAR(10),  
    ShapeID VARCHAR(40),  
    PRIMARY KEY(BusID),  
    FOREIGN KEY(LineID) REFERENCES BusLine(LineID) ON DELETE SET NULL  
);
```

```
CREATE TABLE Users (  
    Username VARCHAR(20) NOT NULL,  
    Password VARCHAR(20) NOT NULL,  
    PRIMARY KEY(Username)
```

);

```
CREATE TABLE FavoriteRoutes (  
  LineID VARCHAR(40) NOT NULL,  
  Username VARCHAR(20) NOT NULL,  
  DateCreated VARCHAR(10),  
  PRIMARY KEY(LineID, Username),  
  FOREIGN KEY(LineID) REFERENCES BusLine(LineID) ON DELETE CASCADE,  
  FOREIGN KEY(Username) REFERENCES Users(Username) ON DELETE CASCADE  
);
```

```
CREATE TABLE FavoriteStops (  
  StopID VARCHAR(20) NOT NULL,  
  Username VARCHAR(20) NOT NULL,  
  DateCreated VARCHAR(10),  
  PRIMARY KEY(StopID, Username),  
  FOREIGN KEY(StopID) REFERENCES BusStop(StopID) ON DELETE CASCADE,  
  FOREIGN KEY(Username) REFERENCES Users(Username) ON DELETE CASCADE  
);
```

Users - 1000 entries:

```
mysql> SELECT COUNT(*) FROM Users;  
+-----+  
| COUNT(*) |  
+-----+  
|      1000 |  
+-----+  
1 row in set (0.02 sec)
```

FavoriteRoutes - 2960 entries:

```
mysql> SELECT COUNT(*) FROM FavoriteRoutes;  
+-----+  
| COUNT(*) |  
+-----+  
|      2960 |  
+-----+  
1 row in set (0.01 sec)
```

FavoriteStops - 2996 entries:

```
mysql> SELECT COUNT(*) FROM FavoriteStops;
+-----+
| COUNT(*) |
+-----+
|      2996 |
+-----+
1 row in set (0.01 sec)
```

BusStop - 1069 entries:

```
mysql> SELECT COUNT(*) FROM BusStop;
+-----+
| COUNT(*) |
+-----+
|      1069 |
+-----+
1 row in set (0.07 sec)
```

Advanced Query 1:

Retrieve all buses that are part of a user's favorite routes, as well as starting at one of their favorite stops

```
SELECT u.Username, r.LineID, s.StopID, b.BusID, COUNT(b.BusID) AS UsefulBusCount
FROM Users u NATURAL JOIN FavoriteRoutes r NATURAL JOIN FavoriteStops s JOIN Buses
b ON r.LineID = b.LineID OR s.StopID = b.StopOrigin
GROUP BY u.Username, b.BusID, r.LineID, s.StopID
ORDER BY u.Username, b.BusID;
```


Second Index added onto Buses-StopOrigin and StopDestination
CREATE INDEX idx_stop on Buses (StopOrign, StopDestination);

Third Index added onto Buses-StopOrigin and StopDestination as well as FavoriteRoutes-DateCreated and FavoriteStops-DateCreated

We will choose the third indexing design, as it provides the same speed up as the second indexing design, where the time it takes to execute the query goes from 0.02 seconds to 0.01 seconds. It combines this speed up along with the benefit of the first indexing design, where it changes the table scan on our FavoriteRoutes table into an index scan, reducing the time of the scan. The index on the stop origin and stop destination greatly help, as it causes the cost of the filter, one of the greatest costing attributes at 2546.95 to be reduced down to a cost of 0, with the time of it being reduced from 12.2 to 0.021. Furthermore, the first inner join's time is also reduced from 12.02 to 12.13, to 2.89 to 11.47, causing it to be much more efficient.

Advanced Query 2:

Retrieve all routes that are currently being utilized for each user, the information corresponding to that route, and the total number of buses that are following that route.

```
SELECT Username, LineID, Number, Direction, COUNT(LineID) AS NumberOfBuses
FROM Buses b NATURAL JOIN BusLine bl NATURAL JOIN FavoriteRoutes f NATURAL JOIN
Users u
GROUP BY Username, LineID, Direction
ORDER BY Username, Number, Direction;
```

Username	LineID	Number	Direction	NumberOfBuses
abjectIguana3	GREY EVENING	70	East	1
abjectIguana3	GREY EVENING	70	West	2
abjectIguana3	YELLOW EVENING	100	North	2
abjectIguana3	YELLOW EVENING	100	South	2
abjectLemur4	GREY EVENING	70	East	1
abjectLemur4	GREY EVENING	70	West	2
abjectLemur4	ILLINI LIMITED EVENING	220	North	1
abjectLemur4	ILLINI LIMITED EVENING	220	South	1
abjectWigeon1	TEAL EVENING	120	East	1
abjectWigeon1	TEAL EVENING	120	West	1
adoringCow3	GREENHOPPER EVENING	50	East	2
affectedQuiche5	GREY EVENING	70	East	1
affectedQuiche5	GREY EVENING	70	West	2
alertPaella9	RUBY EVENING	110	South	1
alertTermite6	SILVER LIMITED EVENING	130	South	1

15 rows in set (0.01 sec)

Original Explain Analyze:

First index was added on Buses-Shape
CREATE INDEX idx_shape on Buses(ShapeID);

Second index was added on Buses-Direction:
CREATE INDEX idx_dir on Buses (Direction);

Third index was added on BusLine-Number and Color
CREATE INDEX idx_num_color on BusLine (Number, Color);

[illegible]

We chose the third indexing design, where we add an index to both the BusLine's number and color. This is because it creates an increase in speed from 0.01 to 0.00 seconds from the original. Although minimal, any increase in speed is good. Because both the first and third indexing design have this increase in speed, we look further into the analysis. The highest cost in both designs is the first nested loop inner join. However, in the first design, we see that the time it takes is between 0.053 and 0.972, however, in the third design, we see that the actual time is 0.037 and 0.946. As such, the third design could be a bit faster than the first. A very apparent change is also seen in the time of the table scan on b, where the time in the third design is 0.013 to 0.017, but in the first, it is 0.023 to 0.028. This is a very drastic difference, nearly double the speed. As a result, we choose the third indexing design due to these increases in speed. It helps that a bus line's number and color is very correlated to its Line ID, and thus is efficient to find.